

# End-to-End Inventory Prediction and Contract Allocation for Guaranteed Delivery Advertising

Wuyang Mao<sup>1</sup>, Chuanren Liu<sup>2\*</sup>, Yundu Huang<sup>1</sup>, Zhonglin Zu<sup>1</sup>, M Harshvardhan<sup>2</sup>, Liang Wang<sup>1</sup>, Bo Zheng<sup>1</sup>

Presented by: M Harshvardhan

<sup>1</sup>Alibaba Group, China

<sup>2</sup>The University of Tennessee, Knoxville

# What is Guaranteed Delivery Advertisement?

## Background of Online Advertisement

→ **Guaranteed Delivery (GD) Advertising** is a strategic approach where advertisers secure their desired inventory of advertising impressions in advance by signing contracts with publishers weeks or months ahead of the targeting dates

→ **Key features of GD Advertising:**

- ◆ Fixed price for impressions
- ◆ Defined target audience
- ◆ Specified target areas
- ◆ Chosen target channels

→ **Advantages of GD over Real-time Bidding**

- ◆ Compared to the dynamic and competitive nature of RTB, GD offers a cost-effective solution for brand advertisers
- ◆ GD ensures wide reach to potential consumers with guaranteed impressions, providing better control over advertising campaigns and budget



## Advertising Allocation Strategies

- Guaranteed Delivery (GD)
- Real-time Bidding (RTB)

# Background and Contributions


---

How does Guaranteed Delivery Allocation work, how does our end-to-end model work, and what are our contributions?

# We visualize End-to-End Prediction and Allocation System

System Architecture for Allocating Ads at Alibaba

- **Our focus is Selling Systems.** The goal is to establish contracts with advertisers in advance by predicting and allocating inventory accurately
- We sign contracts with advertisers in advance while having limited impressions inventory
- The objectives are to:
  - ◆ Maximize inventory sales
  - ◆ Prevent overselling of inventory
- **Online Serving System** ensures
  - ◆ Fulfillment of reserved contracts
  - ◆ Click-Through Rate (CTR) Optimization

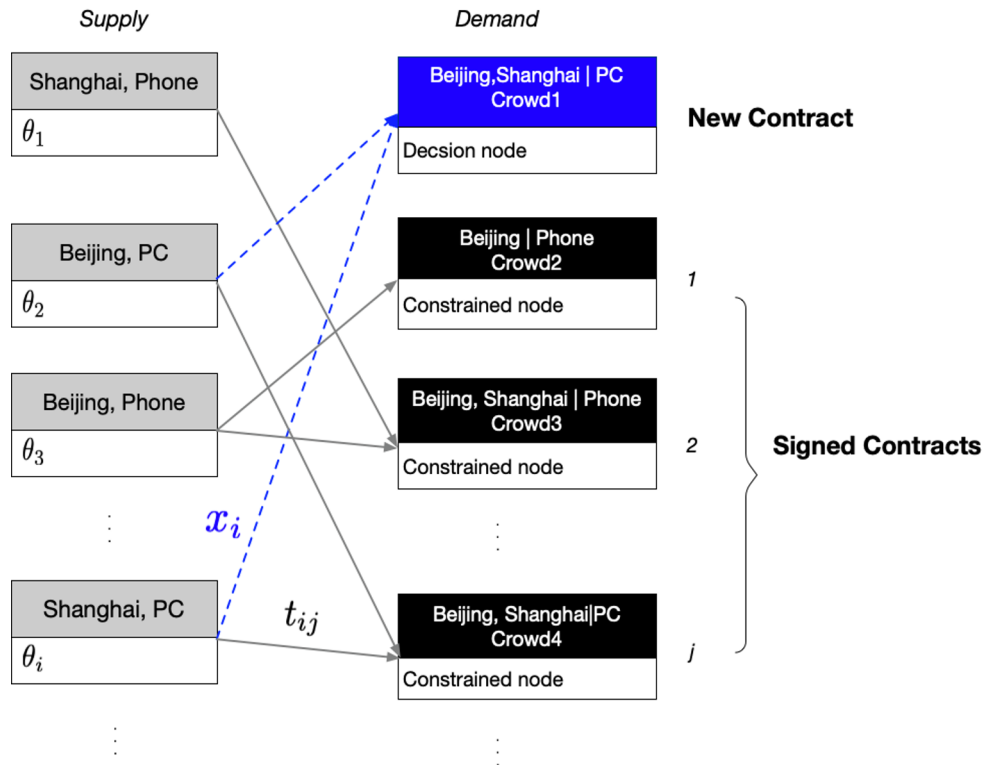


**The system architecture for Guaranteed Delivery (GD) advertising.**

Our model supports the selling system, when advertisers and the publisher sign new GD contracts.

# Traditional methods focus on real-time dispatching of traffic to ads

## Overview of Traditional Two-stage Systems



**Figure:** Bipartite graph of contract allocation problem. Supply nodes are impression inventories while demand nodes are impression contracts.

### Real-Time Bidding (RTB):

An auction-based system where ad impressions are bought and sold on a per-case basis. Advertisers bid for an impression, and if they win, their ad is instantly displayed.

- **Pros:** Dynamic, allows for real-time customization and targeting
- **Cons:** Uncertain costs, less control over ad placement

### Traditional Guaranteed Delivery (GD) Systems:

Advertisers sign contracts with publishers in advance to secure a fixed inventory of ad impressions.

- **Pros:** Secured impressions, cost control
- **Cons:** Less flexibility, may lead to underutilization or overselling of inventory

Given that high-quality impressions often sell out swiftly in large-scale GD marketplaces, even a marginal improvement in inventory usage can significantly increase publisher revenue and serve more advertisers.

# Closing the Gap: Two-Stage Systems vs. Our End-to-End Model

What are the limitations and how we bridge it?

## Limitations of Two-Stage Systems

### Lack of Real-Time Support

- Traditional studies on GD advertising don't provide real-time predictive allocations required for advance contract signing

### Overselling Risk

- Theoretical models for online serving risk overselling if the signed target impressions exceed the inventory limit

### Forecasting Uncertainty

- Two-stage methods assume lower forecasting error translates to better allocation quality, which isn't always the case

### Inefficient Handling of Constraints

- Existing end-to-end learning-based optimization frameworks struggle with large numbers of constraints in GD selling.

## Neural Lagrangian Selling (NLS)

### End-to-End Approach

- We propose an end-to-end approach that integrates inventory prediction and contract allocation, overcoming the limitations of two-stage methods

### Efficient Solver with Less Memory

- Our GD selling system includes an efficient differentiable Lagrangian solver with less memory cost

### Dynamic Feature Extraction

- We use Graph Convolutional Networks (GCNs) to capture dynamic features from advertising contracts, enabling NLS to fit dynamic optimization objectives and constraints

### Improved Performance

- Our approach provides improved allocation error, demonstrating its effectiveness in real-world GD selling systems

# Methodology

---

Deep-dive into how our End-to-End Allocation optimization model works



THE UNIVERSITY OF  
TENNESSEE  
KNOXVILLE

# End-to-End Model Minimizes Allocation Regret Instead of Prediction Loss

## Stage 1: Impression Inventory

- Predictive model (like NN) to forecast impressions
- Let  $\theta$  be the impression inventory of supply nodes;  $z$  has historical and contextual information
- Predict impressions by minimizing loss

$$Loss = ||\theta - g(z; w)||^2$$

## Stage 2: Inventory Allocation

- $\theta$  is supply inventory,  $t$  represents the elements of allocation matrix,  $p$  are penalty constraints (importance of different constraints),  $u$  are slack variables
- $\max_{\{x, t, u\}} (I \odot \theta)'x - p'u,$
- Subject to constraints:

$$\begin{aligned} Gx + Bt - u &\leq h, \\ u &\geq 0, \\ x, t &\in [0, 1] \end{aligned}$$

## Our End-to-End Approach

- NLS minimizes allocation regret, not just prediction loss
- Learns from historical decision cases, considering both inventory prediction and allocation

## Allocation Regret

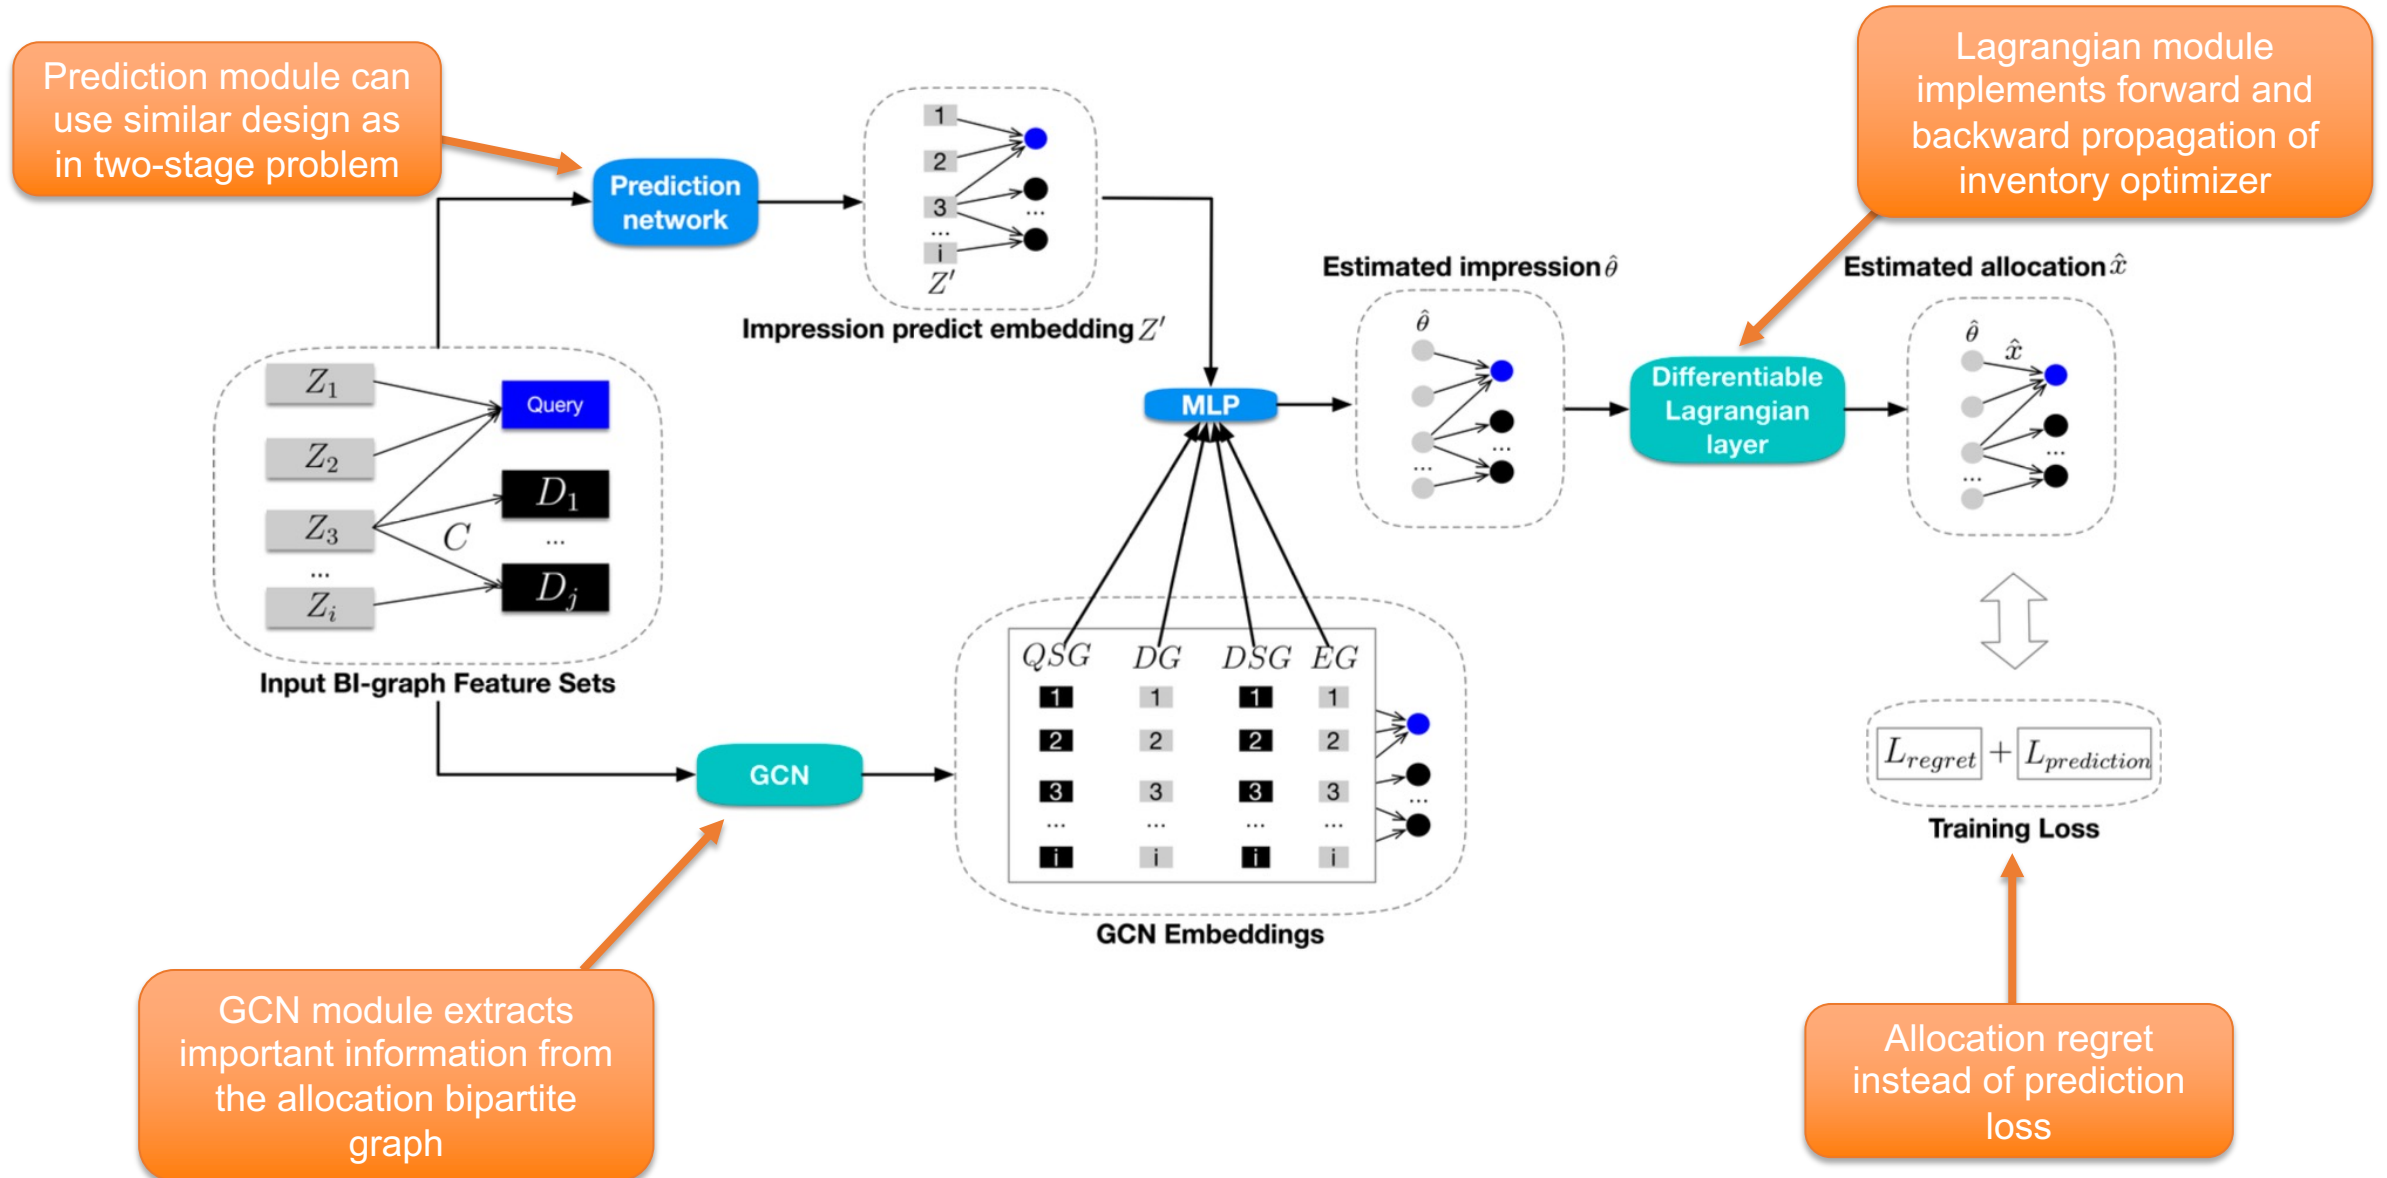
- Formally defined as  $\text{Regret} = \|(I \odot \hat{\theta})^\top \hat{x} - (I \odot \theta)^\top x\|_2^2$
- $\hat{\theta} = g(z; w)$  represents predicted inventories,  $\hat{x}$  corresponds to allocation decisions

## Gradient of Regret

- Gradient of the regret can be computed by differentiating the allocation optimization problem
- Challenge lies in backpropagating optimum  $\frac{\partial \hat{x}}{\partial \hat{\theta}}$



# Architecture of Neural Lagrangian Selling (NLS) Model



# Lagrangian Dual Optimization

- We formulate dual problem to minimize

$$\min_{\alpha, \beta \geq 0} \max_{x, t, u} L(x, t, u, \alpha, \beta),$$

- where the Lagrangian function is

$$L(x, t, u, \alpha, \beta) = \ell(x, t, u) + \alpha^\top (h + u - Gx - Bt) + \beta^\top u.$$

- We derive corresponding KKT (Karush-Kuhn-Tucker) conditions:

$$x = \min(\max(\frac{I \odot \theta - G^\top \alpha}{\lambda}, 0), 1),$$

$$t = \min(\max(\frac{-B^\top \alpha}{\lambda}, 0), 1),$$

$$\alpha \odot (h + u - Gx - Bt) = 0,$$

$$\beta \odot u = 0,$$

$$\alpha + \beta = p.$$

- When  $u > 0$ , then  $\beta = 0$  and  $\alpha = p$ .
- When  $u = 0$ , the gradient is obtained as:

$$\begin{aligned} \frac{\partial L}{\partial \alpha} = & (I \odot \theta - \lambda x - G^\top \alpha) \frac{\partial x}{\partial \alpha} - (B^\top \alpha + \lambda t) \frac{\partial t}{\partial \alpha} \\ & + (h - Gx - Bt), \end{aligned}$$

- We choose hyperparameter  $\lambda = \min(\theta)$
- Small  $\lambda$  results in instability while large  $\lambda$  introduces errors

# Efficient Lagrangian Dual Layer

Integrating the Lagrangian Dual Optimization Layer in Deep Neural Networks

## Solving the GD Selling Allocation Problem

- We integrate the Lagrangian dual optimization layer into predictive deep neural networks, as outlined in *Algorithm 1*.
- Adam gradient descent method is used in Stage 3, with  $b_1, b_2, \mu$  and  $decay\_rate$  as hyperparameters

## Advantages of Our Approach

- **Flexibility:** Our Lagrangian dual layer can be seamlessly embedded into any neural network structure.
- **Autonomy:** Back-propagation can be performed automatically by deep learning frameworks like TensorFlow or Torch.
- **Efficiency:** The computation of the optimization layer avoids complex operations like matrix inversion, enabling easy implementation with batched input and parallel computing for enhanced training/inference efficiency.

---

### Algorithm 1 Forward Pass for Lagrangian Dual Layer (LDL)

---

**Input:**  $\theta$

**Output:**  $x = LDL(\theta)$

Constants:  $G, B, h, p$

Initialization:  $\alpha = 0, mt = 0, vt = 0, \mu = 100, b_1 = 0.9, b_2 = 0.99, e = 10^{-9}, decay\_rate = 0.95$

1:  $\lambda = \min(\theta)$

2: **repeat**

3:   **Stage 1: Calculate original variable**

4:    $x = \frac{I \odot \theta - G^T \alpha}{\lambda}$

5:    $t = \frac{-B^T \alpha}{\lambda}$

6:    $x = \min(0, \max(x, 1))$

7:    $t = \min(0, \max(t, 1))$

8:   **Stage 2: Calculate dual variable gradient**

9:    $\frac{\partial L}{\partial \alpha} = h - Gx - Bt$

10:   **Stage 3: Update dual variable**

11:    $mt = b_1 * mt + (1 - b_1) * \frac{\partial L}{\partial \alpha}$

12:    $vt = b_2 * vt + (1 - b_2) * (\frac{\partial L}{\partial \alpha} \odot \frac{\partial L}{\partial \alpha})$

13:    $\alpha = \min(\max(0, \alpha - \mu * \frac{mt}{\sqrt{vt+e}}, p))$


14:    $\mu = \mu * decay\_rate$

15: **until** Convergence or maximum iterations

---

# Graph Convolutional Network (GCN) Extracts Meaningful Features

- GCN Module extracts meaningful features from the selling allocation problem as objective and conditions vary greatly between GD contracts
- It captures adaptive features for diverse decision scenarios and improve accuracy and flexibility of allocation optimization



Query to Supply GCN	Demand to Supply GCN	Demand GCN	Edge GCN	Concatenate with ReLU
<ul style="list-style-type: none"><li>• Impressions supply is converted to embedding and broadcasted to all supply nodes</li></ul>	<ul style="list-style-type: none"><li>• Aggregated supply features from constraint demand nodes</li></ul>	<ul style="list-style-type: none"><li>• Aggregate demand node to represent competition for supply impressions</li></ul>	<ul style="list-style-type: none"><li>• Edge conditions account for boundary conditions like frequency and crowd</li></ul>	<ul style="list-style-type: none"><li>• All features are concatenated into a single layer passed to Multi-layered Perceptron</li></ul>

# Forward Computation Completes End-to-End Allocation Optimization

Dense Neural Network (DNN) extracts important features from input bi-graph,  $Z'$

$Z'$  along with GCN Embeddings are inputs to Multi-layered Perceptron (MLP) with ReLU activation to estimate supply impressions  $\hat{\theta}$

Estimated supply impressions  $\hat{\theta}$  are input to Langrangian dual layer to produce estimated allocation ratio  $\hat{x}$

Allocation regret is minimized via custom training loss function

$$\left| (I \odot \hat{\theta})' \hat{x} - (I \odot \theta)' x \right|_2^2 + \epsilon \|\theta - \hat{\theta}\|_2^2$$

Training loss has two parts: end-to-end allocation regret and prediction error for inventory nodes. Hyperparameter  $\epsilon$  balances their importances

---

## Algorithm 2 Forward Pass for End-to-End Allocation Optimization

---

**Input:**  $Z^0, D, C$

**Output:**  $x$

*Constant :*  $G, B, h, p, A, I$

1:  $Z' = DNN(Z^0)$

2: **for**  $l = 1, \dots, L$  **do**

3:  $QSG^l = Broadcast(Relu(\frac{I^T Z^{(l-1)} W_{QSG}^l}{\sum_{i=0}^m A_{ij}}))$

4:  $DSG^l = Relu(\frac{A}{\sum_{j=0}^n A_{ij}} \frac{A^T Z^{(l-1)}}{\sum_{i=0}^m A_{ij}} W_{DSG}^l)$

5:  $DG^l = Relu(\frac{ADW_{DG}^l}{\sum_{j=0}^n A_{ij}})$

6:  $EG^l = Relu(\frac{\sum_{j=0}^n (A_{ij} \cdot C_{ij})}{\sum_{j=0}^n A_{ij}} W_{EG}^l)$

7:  $Z^l = Relu((QSG^l, DSG^l, DG^l, EG^l) \cdot W_Z)$

8: **end for**

9:  $\hat{\theta} = Relu(MLP(Z^L, Z'))$

10:  $x = LDL(\hat{\theta})$

---

# Evaluation

---

How does our model perform in comparison to existing methods?

# We evaluate our model via two daily datasets

## Offline Dataset

- The aim is to test the feasibility of end-to-end optimization theory, comparing the two-stage method and end-to-end approach across three inventory allocation cases: full targeting, single targeting, and random targeting
- Basic constraints (overselling and underdelivery) are added to the offline inventory allocation cases
- *Supply x Demand* dimension is  $100 \times 10$  with only 110 constraints (100 overselling and 10 underdelivery)

## Online Dataset

- Experiments are conducted on two online advertising selling datasets: Pre-Video Ads (PVA) and Open-Screen Ads (OSA)
- Each supply node of PVA represents a *Channel x City x Position* combination; for OSA is *City x App*
- The inventory allocation and allocation problems are more dynamic and complex compared to the offline datasets
- For  $m$  supply nodes and  $n$  demand nodes, there are  $(m \times n + m + n)$  constraints

Dimensions	Offline			Online PVA	Online OSA
	full targeting	single targeting	random targeting	real targeting	real targeting
<i>supply <math>\times</math> demand</i>	$100 \times 10$	$100 \times 10$	$100 \times 10$	$500 \times 20$	$500 \times 50$
<i>constraints</i>	110	110	110	<b>10520</b>	<b>25550</b>

# We compare our end-to-end model with five benchmark models on six metrics

## Benchmark Models

### Two-stage Model

- Compares end-to-end approach with traditional two-stage method

### Pure Fully-Connected (PF)

- Baseline end-to-end approach using a simple black-box neural network

### Pure Prediction GCN (PPG)

- Removes Lagrangian layer from NLS, uses GCN for prediction

### Prediction Network + Lagrangian solver (PL)

- End-to-end approach without GCN module

### GCN+QPTL/GCN+InOpt (End-to-End)

- Implements established LP solvers IntOpt and QPTL for comparison with Lagrangian layer

## Evaluation Metrics



End-to-End Normalized Deviation Error



First-stage Error



Second-stage Error



Publisher Revenue (Avg. Revenue per Day)



Delivery Rate (Delivered/Promised)



Usage Rate (Sold/Available Impressions)




# NLS outperforms all other models on most benchmarks

Methods	full targeting		single targeting		random targeting	
	$ND_{pre}$	$ND_{reg}$	$ND_{pre}$	$ND_{reg}$	$ND_{pre}$	$ND_{reg}$
Two Stage	$0.101 \pm 0.003$	$0.023 \pm 2e-4$	$0.101 \pm 0.003$	$0.125 \pm 0.005$	$0.101 \pm 0.003$	$0.045 \pm 0.002$
PF	$0.130 \pm 0.010$	$0.045 \pm 0.005$	$0.115 \pm 0.008$	$0.132 \pm 0.010$	$0.121 \pm 0.010$	$0.076 \pm 0.007$
PPG	$0.125 \pm 0.010$	$0.015 \pm 1e-4$	$0.127 \pm 0.007$	$0.112 \pm 0.001$	$0.135 \pm 0.011$	$0.036 \pm 0.001$
PL	$0.102 \pm 0.001$	$0.008 \pm 1e-4$	$0.103 \pm 0.001$	$0.113 \pm 0.003$	$0.101 \pm 0.002$	$0.047 \pm 2e-4$
<b>NLS</b>	<b><math>0.096 \pm 0.002</math></b>	<b><math>0.007 \pm 2e-4</math></b>	<b><math>0.097 \pm 0.001</math></b>	<b><math>0.098 \pm 0.001</math></b>	<b><math>0.095 \pm 0.001</math></b>	<b><math>0.029 \pm 1e-4</math></b>

Table: Experiment Results on Offline Data

Methods	PVA		OSA	
	$ND_{pre}$	$ND_{reg}$	$ND_{pre}$	$ND_{reg}$
Two Stage	$0.069 \pm 0.002$	$0.068 \pm 0.004$	$0.067 \pm 0.002$	$0.132 \pm 0.005$
PF	$0.083 \pm 0.015$	$0.095 \pm 0.020$	$0.075 \pm 0.015$	$0.128 \pm 0.021$
PPG	$0.085 \pm 0.005$	$0.054 \pm 0.002$	$0.078 \pm 0.003$	$0.086 \pm 0.004$
PL	$0.065 \pm 0.002$	$0.061 \pm 0.002$	<b><math>0.065 \pm 0.001</math></b>	$0.136 \pm 0.004$
<b>NLS</b>	<b><math>0.064 \pm 0.003</math></b>	<b><math>0.041 \pm 0.001</math></b>	$0.068 \pm 0.003$	<b><math>0.058 \pm 0.001</math></b>

Table: Experiment Results on Online Data



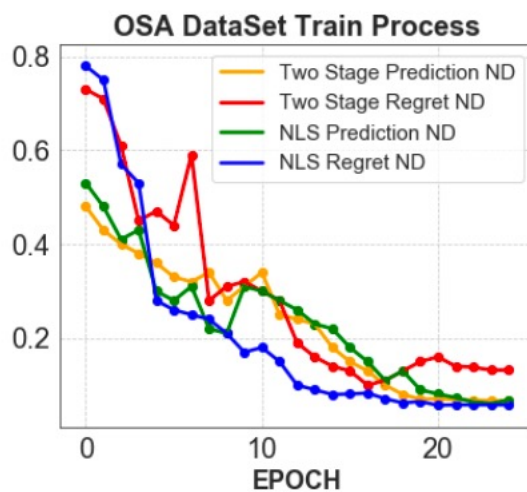
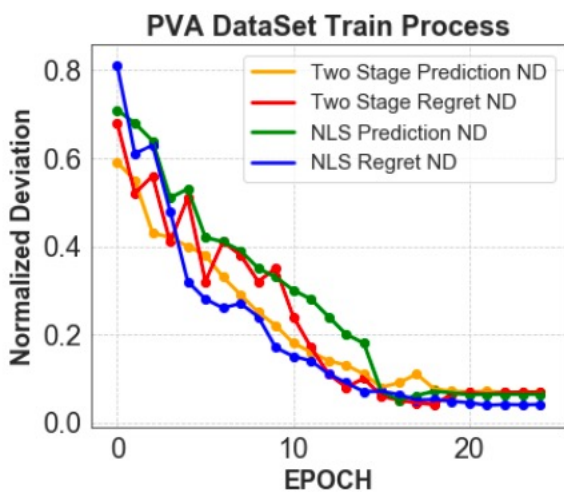
Our NLS has fewer outliers in comparison with two-stage methods.

$ND_{pre}$  and  $ND_{reg}$  are Normalized Deviations for prediction and regret.

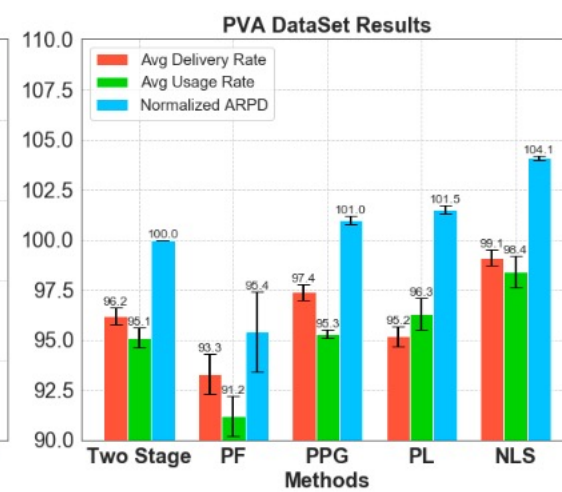
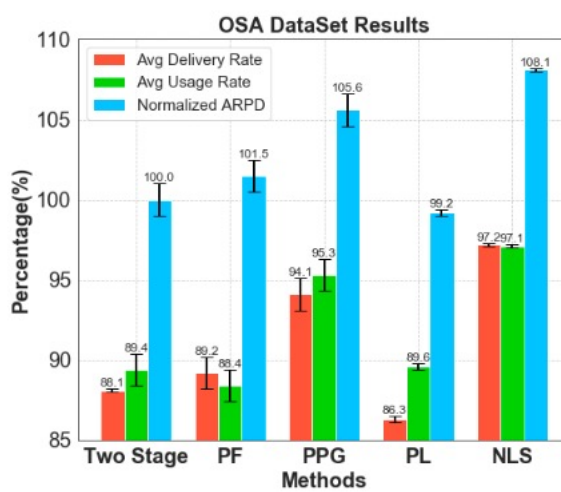
# NLS handles complex optimization constraints better than benchmarks

Methods	Offline (random)			PVA			OSA		
	$ND_{pre}$	$ND_{allo}$	$ND_{reg}$	$ND_{pre}$	$ND_{allo}$	$ND_{reg}$	$ND_{pre}$	$ND_{allo}$	$ND_{reg}$
GCN + QPTL	$0.098 \pm 0.008$	$0.15 \pm 0.02$	$0.032 \pm 0.005$	-	-	-	-	-	-
GCN + IntOpt	$0.100 \pm 0.003$	$0.25 \pm 0.03$	$0.038 \pm 0.005$	$0.068 \pm 0.006$	$0.32 \pm 0.04$	$0.063 \pm 0.002$	$0.070 \pm 0.008$	$0.33 \pm 0.04$	$0.083 \pm 0.006$
<b>NLS</b>	<b><math>0.095 \pm 0.001</math></b>	<b><math>1e-4 \pm 1e-5</math></b>	<b><math>0.029 \pm 1e-4</math></b>	<b><math>0.064 \pm 0.003</math></b>	<b><math>7e-4 \pm 1e-5</math></b>	<b><math>0.041 \pm 0.0031</math></b>	<b><math>0.068 \pm 0.003</math></b>	<b><math>6e-4 \pm 1e-5</math></b>	<b><math>0.058 \pm 0.001</math></b>

NLS outperforms classic optimization methods on all benchmarks



NLS has more stable training trajectory



NLS has better delivery rate, usage rate and publisher revenue

$ND_{pre}$  and  $ND_{allo}$  are Normalized Deviations for prediction and allocation.  
ARPD is Average Revenue per Day.

# Conclusion: Key Takeaways and Results

1. **Key Idea:** Proposes Neural Lagrangian Selling (NLS), an end-to-end approach for inventory prediction and contract allocation in Guaranteed Delivery (GD) advertising.
2. **Methodology:** NLS integrates a differentiable Lagrangian dual optimization layer into predictive deep neural networks, overcoming limitations of traditional two-stage methods.
3. **Comparison:** NLS is compared with established models and methods, including Two-Stage, PF, PPG, PL, GCN+QPTL, and GCN+InOpt.
4. **Results:** NLS outperforms existing models in terms of allocation regret minimization and computational efficiency.
5. **Significance:** The end-to-end approach offers significant improvements in the GD selling process by better handling dynamic optimization objectives and constraints.

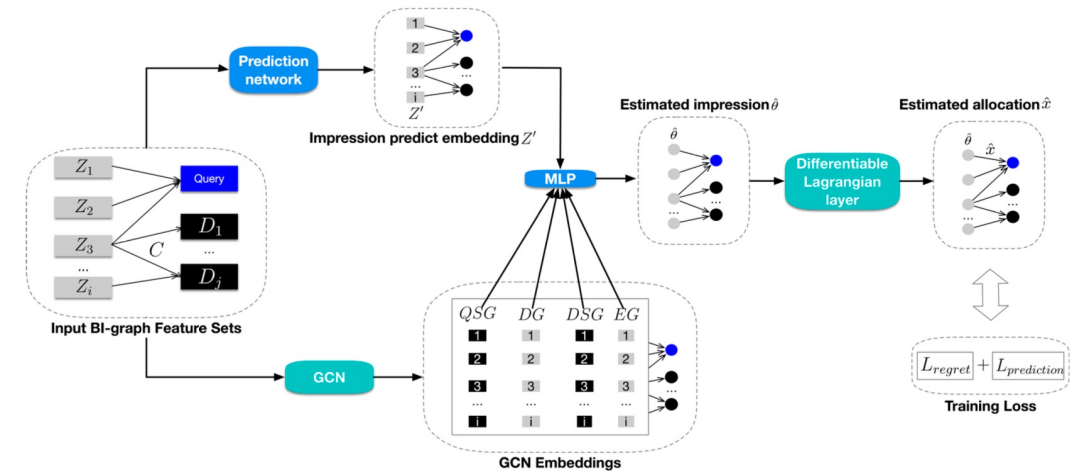


Figure: Neural Lagrangian Selling (NLS) Model for End-to-End Prediction and Optimization

# Thank you for your time.

---

Questions?



THE UNIVERSITY OF  
TENNESSEE  
KNOXVILLE